

---

# Reinforcement Learning

## Assignment 2

### Checkpoint 1

---

**Sumeet Aher**

sumeetmi  
sumeetmi@buffalo.edu

**Suraj Sah**

suraj  
suraj@buffalo.edu

**"I certify that the code and data in this assignment were generated independently, using only the tools and resources defined in the course and that I did not receive any external help, coaching or contributions during the production of this work."**

#### Abstract

- Explore 'CartPole-v1' and provide the main details about the environment (e.g. possible actions/states, goal, rewards, etc).
- Choose one more environment that you use for the assignment and provide the main details about it

#### 1 Exploring CartPole-v1

A pole is attached by an un-actuated joint to a cart, which moves along a frictionless track. The system is controlled by applying a force of +1 or -1 to the cart.

##### 1.1 Action Space

`Discrete(2)`

The Discrete space allows a fixed range of non-negative numbers, so in this case valid actions are either 0 or 1 which represents application of the force +1 or -1.

##### 1.2 Observation Space

`Box(4, )`

The Box space represents an n-dimensional box, so valid observations will be an array of 4 numbers.

Observation Space bounds

33 Observation\_space.high: [4.8000002e+00 3.4028235e+38  
34 4.1887903e-01 3.4028235e+38]  
35 Observation\_space.low: [-4.8000002e+00 -3.4028235e+38  
36 -4.1887903e-01 -3.4028235e+38]

37

### 38 1.3 Goal

39 The goal for this environment per episode is to keep the pendulum upright. So the  
40 pendulum keeps falling and in the effort to keep this pendulum upright, the cart  
41 may move on the axis. For keeping the episode end in check, there are two  
42 checks. Either the pole is more than 15degrees from vertical or the cart moves  
43 more than 2.4 units from the centre.

44

### 45 1.3 Reward

46 A reward of +1 is provided for every timestep that the pole remains upright

47

## 48 2 Exploring LunarLander-v2

49

50 The Lunar Lander will always land on (0,0), so the landing coordinates are fix.

51

### 52 2.1 Action Space

53 Discrete(4)

54 The Discrete space allows a fixed range of non-negative numbers, so in this case  
55 valid actions are 0, 1, 2, 3.

56

### 57 2.2 Observation Space

58 Box(8,)

59 The Box space represents an n-dimensional box, so valid observations will be an  
60 array of 8 numbers.

61 Observation Space bounds-

62 Observation\_space.high: [inf inf inf inf inf inf inf inf]

63

64 Observation\_space.low: [-inf -inf -inf -inf -inf -inf -inf -inf]

65

66

### 67 2.3 Goal

68 The goal for this environment per episode is to land on the coordinates.

69

### 70 2.3 Reward

- 71 • When the lander touches the ground, there is an additional reward of  
72 +100/-100 depending on the coordinates.
- 73 • As soon as any leg touches the ground, it gets a reward of +10
- 74 • There are three actions, from them firing the main engine will cost the  
75 agent -0.3 reward points
- 76 • On succesfull landing, the agent gets a reward of 200 points.

77

### 3 Grid World Used

The environment defined here is of a Kitchen. The agent in this Kitchen is an Ant whose objective is to reach her hole. In the environment grid there are small heap of sugar in form of positive rewards and pesticide in form of negative rewards. The environment has the following properties –

Agent: An Ant

States: 16 states in form of 4\*4 square.

Actions: 4 actions {Right, Left, Up, Down}

Rewards: 5 Rewards in form of sugar and pesticide with values {-10,10,10,10,50} at location {(0,2), (1,1), (1,3), (3,1),(3,3)}

Main Objective: To reach hole home (3,3)

The changes introduced in environment for DQN –

1. The reward for going close +1 and -1 if going away from goal
2. All reward vanishes after it's taken for 1st time

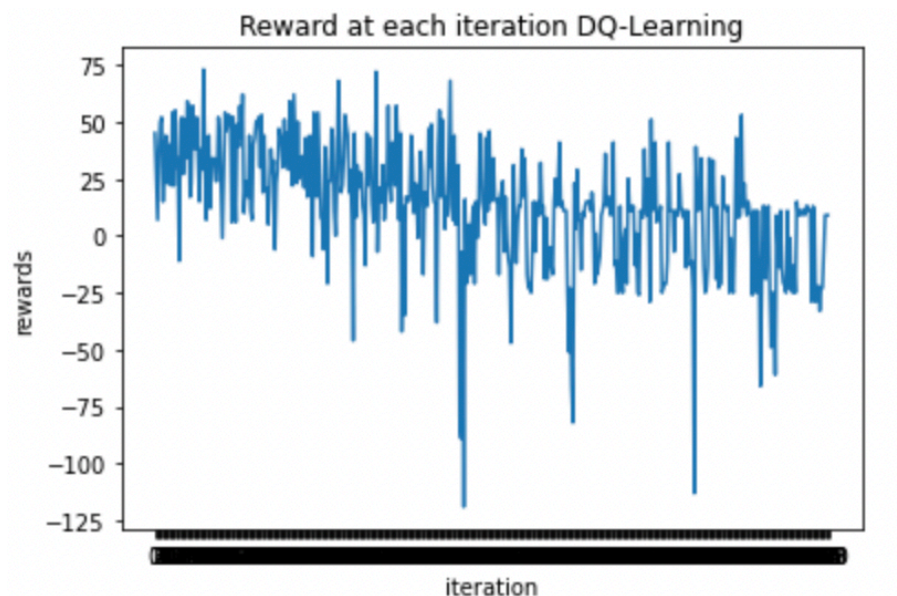


- **Using experience replay in DQN and how its size can influence the results**

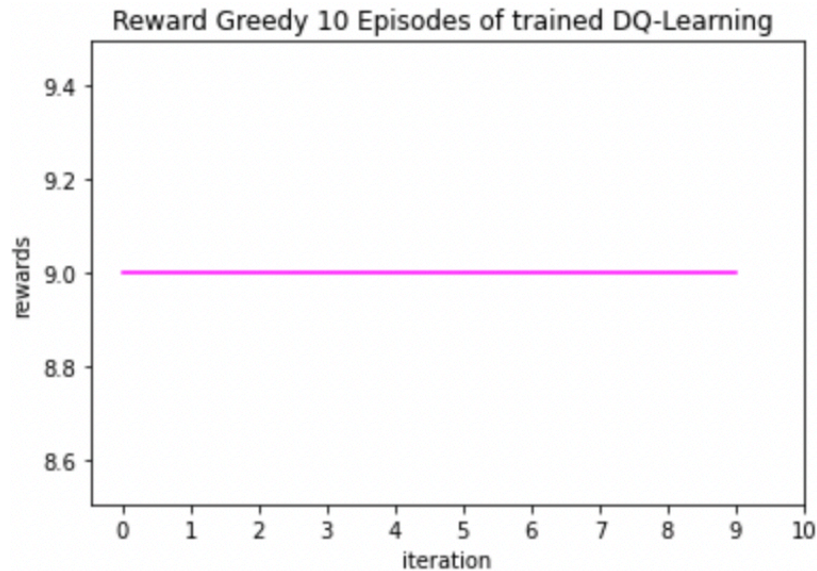
-The experience replay was implemented using a list of 1000 capacity. Which is overwritten with each step tuple from start to end.

The size if increased will need the network to be trained more as the number of samples taken from the replay will slowly be updated across the list. This will result more time in training and more number of episodes for training as the list properties needs to be updated across many steps

- **Introducing the target network**
  - The target network has the same architecture as the policy network. We copy the weights of the policy network to target network after every 5 episodes of training.
  - The target value is calculated from this network which is used as input to the loss function calculation and gradient decent of the policy network
- **Representing the Q function as  $q^*(s, w)$** 
  - The Q function  $q^*(s, w)$  is a three layer neural network with input as flatten and two 128 dimension hidden layer and output with  $Q(s, a)$  value corresponding to each action.



From the cumulative reward graph , we can deduce the training was unstable and didn't increased in value over time



124

125 The final output on 10 episode run on the target network we see  
 126 the reward comes to be 9. This compared to vanilla Qvalue table  
 127 training which gave 31 output shows the trained model is not  
 128 performing that well

## Final Submission

129

130

131 **3 Discuss the algorithm you implemented.**

132 To get the target in Vanilla DQN, we wrote the following:

133  $Y = \text{reward} + \gamma * (1 - \text{done\_}) * \max\_q$

134 Where *done* is the tensor thus keeping only those maximum q values which  
 135 are not going to terminate in the next state.

136 To get the target in the Double DQN, we wrote the following:

```

137 next_q =
138 self.model_policy.predict(tf.reshape(state_next, (minibatch_size, obs_size)))
139 max_a = tf.math.argmax(next_q, 1)
140 double_q =
141 self.model_target.predict(tf.reshape(state_next, (minibatch_size, obs_size)))
142 target_y = reward_ + gamma*(1-done_)*(tf.gather(double_q[0],
143 max_a).numpy())
  
```

146 This code first predicts the action with maximum q value and then using this  
147 action we find the predicted value from the target model.

148 We use this value as the ground truth for back-propagation.


149 Algorithm implemented as an improvement to Vanilla DQN is Double DQN.

#### 150 **4 What is the main improvement over the vanilla DQN?**

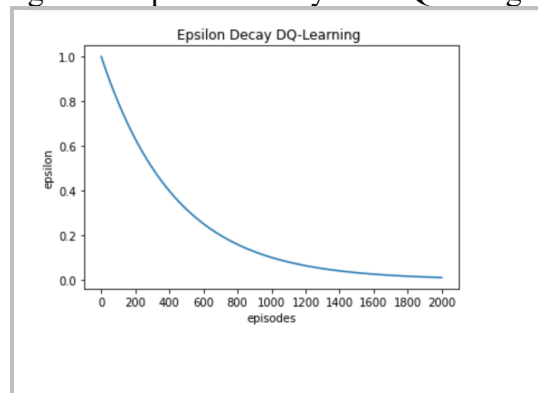
151 The main improvement expected over the vanilla DQN is faster convergence  
152 as well not getting stuck at a local optimum thus not stopping the agent  
153 from moving over to a higher reward.

#### 154 **5 Show and discuss your results after applying your the** 155 **two algorithms implementation on the environment.** 156 **Plots should include epsilon decay and the reward per** 157 **episode.**

158 Below are the results for the grid environment, on application of DQN and  
159 then double DQN:




160 Figure 1: Epsilon Decay for DQN on grid

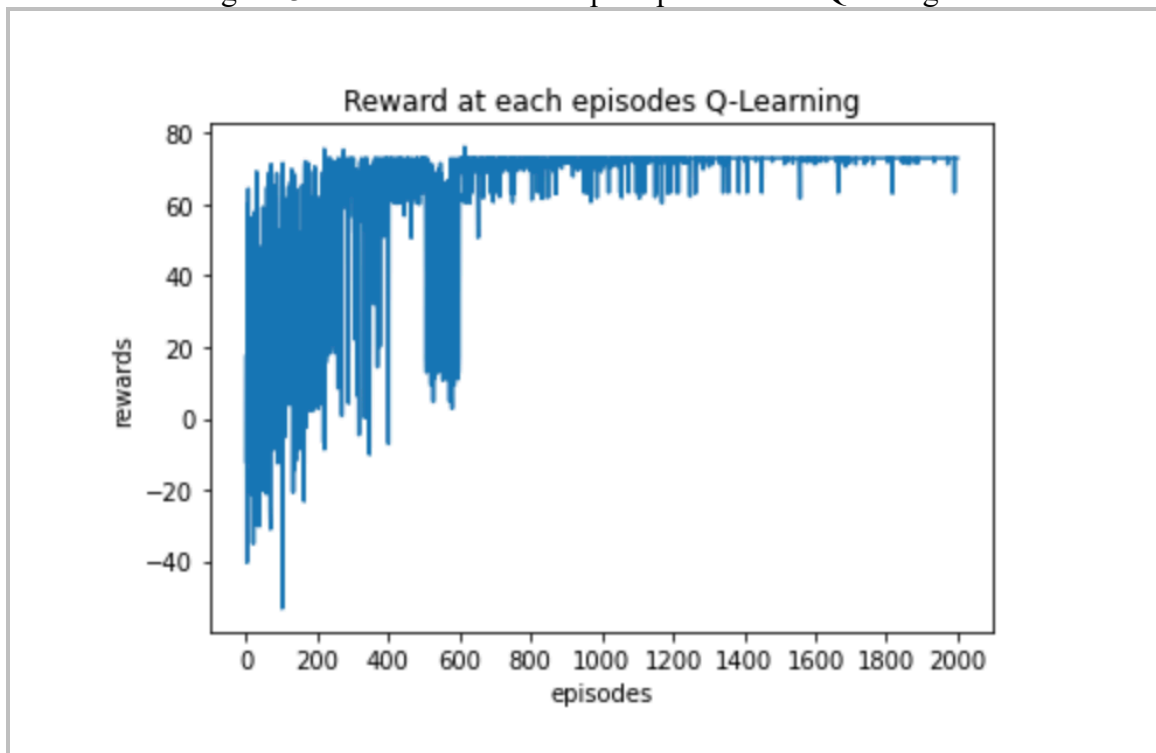


161 Figure 2: Epsilon Decay for Double DQN on grid


162



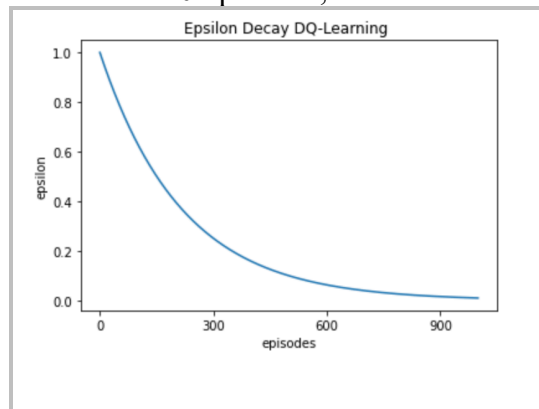
164 Figure 3: Cumulative reward per episode for DQN on grid




165 Figure 4: Cumulative reward per episode for Double DQN on grid



166 Figure 1: Epsilon Decay for DQN on cartpole ( as we get average reward of  
 167 >470 for consecutive 10 episodes, we conclude the training )



168 Figure 2: Epsilon Decay for Double DQN on cartpole

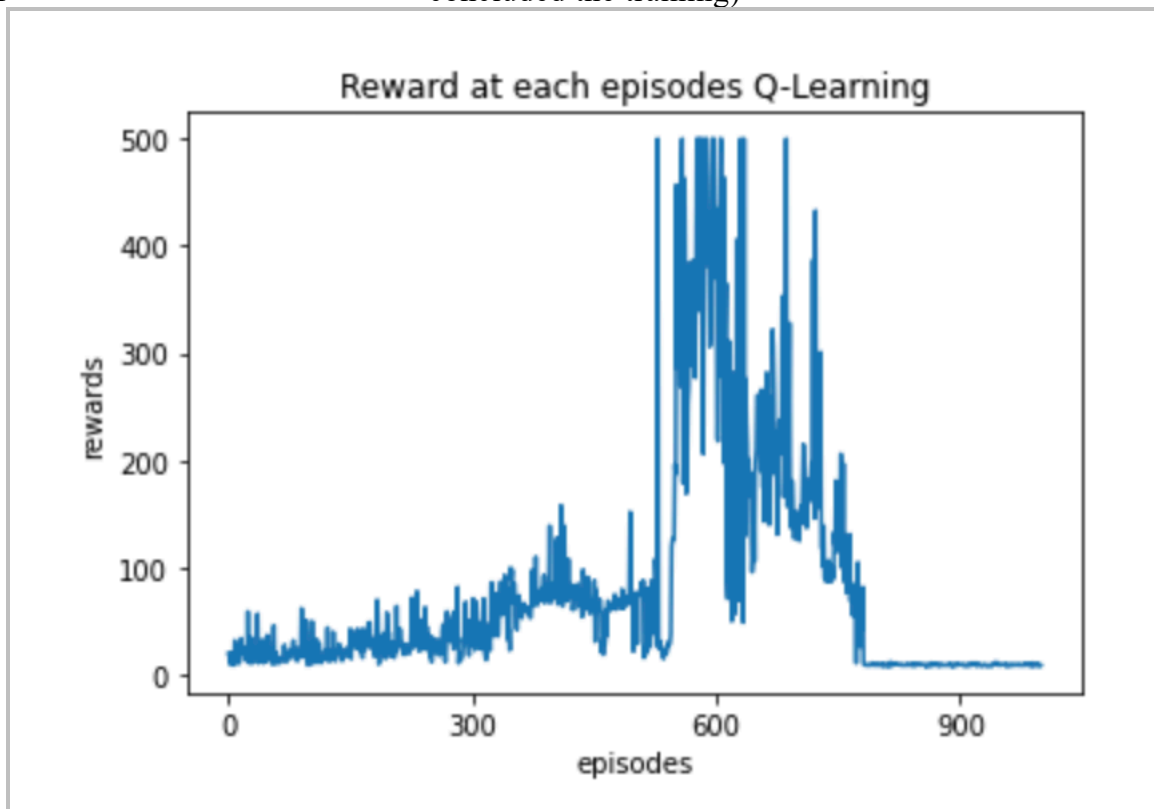


169 Figure 3: Cumulative reward per episode for DQN on cartpole  
 170 (as we got an average reward of >470 for 10 consecutive episodes, we



171

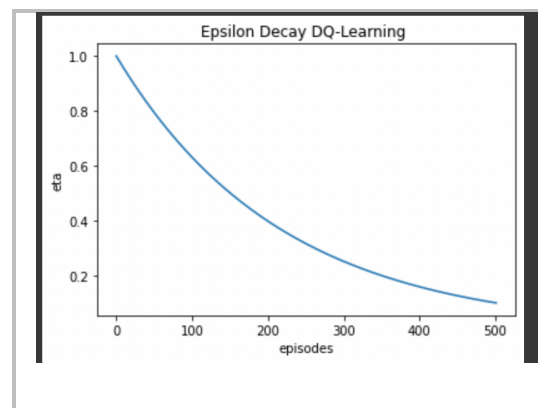
concluded the training)



172

Figure 4: Cumulative reward per episode for Double DQN on cartpole

173



174

Figure 1: Epsilon Decay for DQN on Lunar Lander




Figure 2: Epsilon Decay for Double DQN on Lunar Lander





Figure 3: Cumulative reward per episode for DQN on Lunar Lander




179      Figure 4: Cumulative reward per episode for Double DQN on Lunar Lander

180      1) For Grid environment, Double DQN seems to be more robust  
 181              and on later stages doesn't show any signs of comebacks,  
 182              where as there is fluctuation in normal DQN training in the  
 183              later stage but that is taken care of by the model.

184      2) For CartPoleV1, for just DQN, the model gave a consecutive  
 185              average of >470 in the initial 400 episodes itself, and we  
 186              ended the training there. Further we used the model\_target  
 187              to predict the next 10 episodes of the agent and we got a  
 188              500 score on each of the episode using greedy approach. For  
 189              Double DQN, we seem to have a training issue, as the agent  
 190              collects the maximum reward but then slides down to lower  
 191              rewards. Tuning the hyper parameters will solve this  
 192              problem.


193      3) For Lunar Lander, we had to abruptly end the training around  
 194              550 episodes for both DQN and double DQN, hence the  
 195              epsilon can be seen decayed only upto a certain amount as  
 196              it was supposed to decay to 0.001 by 1000 episodes. But, till  
 197              600 episodes, we can see that the agent was slowly learning  
 198              and having increasing rewards.


199   **6   Provide the evaluation results. Run your environment**  
200   **for at least 5 episodes, where the agent chooses only**  
201   **greedy actions from the learnt policy. Plot should**  
202   **include the total reward per episode.**



203                      Figure 5: DQN on grid world

204






205

Figure 6: Double DQN on grid world

206



207

Figure 7: For DQN on CartPole




Figure 8: For Double DQN on CartPole




Figure 9: For DQN on Lunar Lander




Figure 10: For Double DQN on Lunar Lander

7 **Compare the performance of both algorithms (DQN & Improved version of DQN) on the same environments (e.g. show one graph with two reward dynamics) and provide your interpretation of the results. Overall three rewards dynamics plots with results from two algorithms applied on:**

- **Grid-world environment**




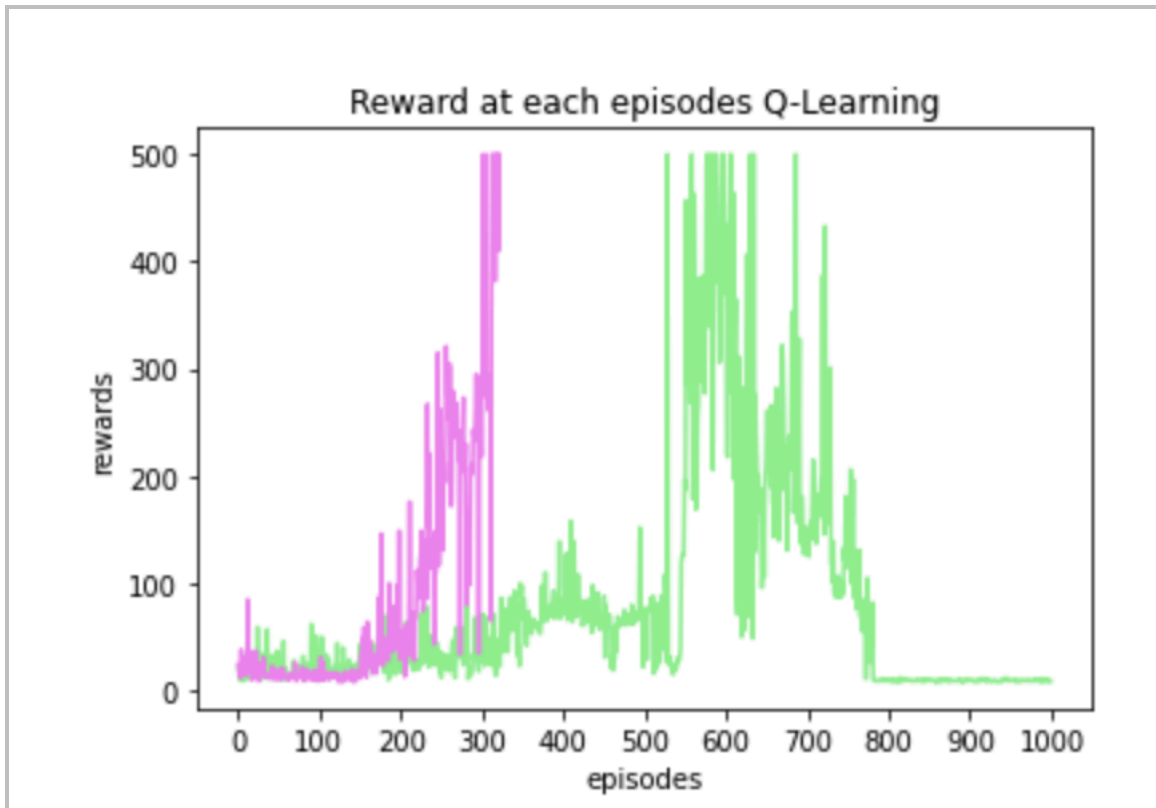
Figure 11: Pink is Rewards per episode for DQN on Grid environment

Figure 12: Yellow is Rewards per episode for Double DQN on Grid Environment

The DQN graph seems to be fluctuating and looks like the model isn't learning properly, on the other hand, the Double DQN seems to be in a better condition with a clear reward exposure in the beginning and a somewhat constant behaviour later.

- 'CartPole-v1'






231                      Figure 13: Pink Reward per episode for DQN on cartpole


232                      Figure 13: Green Reward per episode for Double DQN on cartpole

233      Double DQN seems to have fluctuated in the second half episodes after 500  
 234      where it starting catching big rewards but was brought down either by the  
 235      target Q model or the big Buffer size.


236

237                      • **Lunar Lander**





238 Figure 15: Reward per episode for DQN on Lunar Lander



239 Figure 16: Reward per episode for Double DQN on Lunar Lander

240 For DQN, the model seems to be learning at a good rate and accumulating  
241 rewards. For Double DQN, it seems that the model is fluctuating and would  
242 need some hyperparameter tuning for a slow and robust learning.

243 **8 Provide your interpretation of the results. E.g. how the**  
244 **same algorithm behaves on different environ- ments, or**  
245 **how various algorithms behave on the same environment.**

246 It seems that Double DQN is giving better results by not  
247 localizing too early compared to normal DQN for the grid  
248 environment. DQN seems to need different hyperparameters to  
249 train compared to double DQN for all environments. Double  
250 DQN, which is dependant on the policy model for target  
251 calculation along with target model unlike normal DQN seems to  
252 be very sensitive compared to normal DQN. For Lunar Lander and  
253 CartPole, a better combination of hyperparameters will give us  
254 better results.

255 **Contribution:**

Team Member	Assignment Part	Contribution
Suraj	Part1,2,3	50%
Sumeet	Part1,2,3	50%

256

257 **References**

258 [1] <https://gym.openai.com/docs/>

259 [2] <https://gym.openai.com/envs/LunarLander-v2/>

260 [3] <https://gym.openai.com/envs/CartPole-v0/>

261 [3] <https://web.stanford.edu/class/psych209/Readings/MnihEtAlHassibis15NatureControlDeepRL.pdf>

262 [4] Playing Atari with Deep Reinforcement Learning – Volodymyr eta.

263 [5] [https://medium.datadriveninvestor.com/training-the-lunar-lander-agent-with-deep-q-learning-and-](https://medium.datadriveninvestor.com/training-the-lunar-lander-agent-with-deep-q-learning-and-its-variants-2f7ba63e822c)  
264 [its-variants-2f7ba63e822c](https://medium.datadriveninvestor.com/training-the-lunar-lander-agent-with-deep-q-learning-and-its-variants-2f7ba63e822c)

265